

Operating Systems II

Device Drivers

Device Driver Motivation

- ❑ Multiple processes trying to access the same hardware.
 - Need to ensure integrity of the device.
 - ❑ Quote: “Device drivers are a collection of routines which write magical numbers to magical places in the hardware”.
 - But, how does one integrate these magical routines into the OS?
-

OS Responsibilities

- ☐ When it comes to devices, your OS has several responsibilities:
 - has to provide full access to the features of the hardware device
 - needs to coordinate hardware resources.
 - ☐ Only one task can manipulate a device at a time.
 - needs to protect users from the device.
 - ☐ Incorrect actions by the device driver can crash the system.
-

Driver Overview

- ❑ Device drivers are executed as part of the system kernel.
 - ❑ Drivers are either part of the kernel image, or loaded as a module (statically or dynamically).
 - ❑ In either case it must provide a **standard** interface by which ALL devices can be accessed.
 - Necessary for kernel flexibility.
-

Everything is a File

- ❑ You have probably heard the saying that in UNIX, everything is a file.
 - ❑ This is true, because ALL devices are provided as a file in the UNIX OS.
 - ❑ Special entry points are provided to the driver through virtual files.
 - ❑ The user can then perform basic file operations on the device
 - read, write, open, close, ...
-

Major & Minor Numbers

- ❑ Devices require unique identification in the system
 - Major and Minor numbers?
 - ❑ The major number identifies the type of device
 - ❑ The minor number identifies any mode or subunit of the device.
 - Check out Linux's `Documentation/devices.txt`
-

Blocks & Characters

- Devices within the kernel can be separated between character and block devices.
 - Block devices work on the basis of blocks of memory or the management of block sizes.
 - Character devices work on devices character by character or in sequential byte streams.
-

Example from devices.txt

1 char Memory devices

1 = /dev/mem	Physical memory access
2 = /dev/kmem	Kernel virtual memory access
3 = /dev/null	Null device
4 = /dev/port	I/O port access
5 = /dev/zero	Null byte source
6 = /dev/core	OBSOLETE - replaced by /proc/kcore
7 = /dev/full	Returns ENOSPC on write
8 = /dev/random	Nondeterministic random number gen.
9 = /dev/urandom	Faster, less secure random number gen.
10 = /dev/aio	Asynchronous I/O notification interface

block RAM disk

0 = /dev/ram0	First RAM disk
1 = /dev/ram1	Second RAM disk

...

250 = /dev/initrd	Initial RAM disk {2.6}
-------------------	------------------------

Creating a Virtual Driver File

- ❑ One can create a virtual driver file in Linux by using the `mknod` command.
 - `mknod /dev/name type major minor`
 - `/dev/name` the virtual file for the device.
 - `type` indicates whether it is a block or character device.
 - `major/minor` the unique identifier for device.
-

Device Drivers as a Module

- ❑ If a device is provided as a module, then it must be registered and unregistered with the kernel.
 - ❑ To register a device driver
 - `register_chrdev()` or `register_blkdev()`
 - ❑ To unregister a device driver
 - `unregister_chrdev()` or `unregister_blkdev()`
-

Accessing the Device

- ❑ Once created, the device is accessible through a standard interface.
 - Function pointers in a structure.
 - ❑ Two types of structures typically exist, one for character devices and one for block devices.
-

file_char_op Structure

- ❑ Contains the mapped functions to use when manipulating a character device.
 - ❑ The user will fill the structure with mapped functions and then pass it as a parameter in registering the device.
 - ❑ Internally the kernel keeps this structure in a hash table for character devices that can be quickly accessed by major and minor numbers.
-

`file_char_op` (cont.)

- ❑ `owner` - reference to the registered module.
 - ❑ `llseek` - updates the file pointer.
 - ❑ `read` - read `n` bytes starting at given offset.
 - ❑ `write` - writes `n` bytes starting at given offset.
 - ❑ `readdir` - returns the next directory entry.
 - ❑ `ioctl` - sends a command to underlying hardware device.
-

file_char_op (cont.)

- ❑ `poll` - checks if activity, sleeps until something happens.
 - ❑ `mmap` - maps the file into a process address space.
 - ❑ `open` - creates a file object linked to object.
 - ❑ `close` - closes file linked to object.
 - ❑ `flush` - flushes file (filesystem dependent).
-

file_char_op (cont.)

- ❑ `release` - releases the file object.
 - ❑ `fsync` - writes all cached data of file to disk.
 - ❑ `fasync` - enables/disables asynchronous I/O notification.
 - ❑ `lock` - applies a lock to the file.
 - ❑ `readv` - reads bytes from file and puts data into vectors.
-

file_char_op (cont.)

- ❑ `writew` - writes bytes from vectors and puts into file.
 - ❑ `sendpage` - transfers data from one file to another (used by sockets).
 - ❑ `get_unmapped_area` - gets unused address range to map the file (frame buffer memory mappings).
-

file_block_op Structure

- ❑ Contains the mapped functions to use when manipulating a block device.
 - ❑ The user will fill the structure with mapped functions and then pass it as a parameter in registering the device.
 - ❑ Internally the kernel keeps this structure in a unique hash table for block devices that can be quickly accessed by major and minor numbers.
-

file_block_op Structure

- ❑ open - open the block device file.
 - ❑ release - close the last reference to block device file.
 - ❑ ioctl - issue a i/o control system call on block device.
-

file_block_op Structure

- ☐ `check_media_change` - check if the media has been changed (e.g. floppy drive)
 - ☐ `revalidate` - check if the block device holds valid data.
 - ☐ `owner` - reference to the registered module.
 - ☐ Other basic operations ...
-

Linux Example: `/dev/random`

- ❑ You can check out `devices/char/random.c` for more info on this example.
 - ❑ The random device simply returns a random number to the reading process.
 - ❑ So it only registers the `read`, `write`, `poll`, and `ioctl` functions.
-

Polling and Interrupts

- ❑ Devices are slower than the processor.
 - ❑ Waiting for devices to perform operations is a waste of processing time.
 - ❑ Device drivers typically call schedule to move the processor over to another task.
 - ❑ But how does the processor know the device has completed?
 - Several solutions including interrupts and polling.
-

Polling

- ❑ Certain devices can work under the conditions that the processor can ask the device if it has completed the operation.
 - ❑ In this case, the processor has to regularly check the device for completion.
 - ❑ The task calls `schedule()` to preempt itself, but does not put itself to sleep!
 - The device will be polled when the task is next scheduled.
-

Interrupt Mode

- ❑ In interrupt mode the device informs the processor that it is completed its operation via an IRQ (Interrupt Request).
 - ❑ This relies on the physical hardware device being capable of generating an interrupt.
 - ❑ When an IRQ is generated, the processor switches to execute an interrupt service routine (ISR).
-

Interrupt Mode (cont.)

- ❑ Using interrupts, the device driver operation invokes `interruptible_sleep_on()` putting the task to sleep.
 - ❑ The ISR upon receiving the IRQ, performs any additional work and signals the original task to resume.
 - ❑ Example - mouse.
 - Every movement causes an IRQ
 - ISR reads the port and passes it along to the application.
-

Allocating IRQ's

- ❑ IRQ's have to be requested. Not hard coded to each device.
 - Impossible to predict the different combinations of hardware equipment in a computer.
 - ❑ Allocated through the use of:
 - `int request_irq(unsigned int irq, void (*handler_func)(), unsigned long irqflags, char* devname, void* dev_id)`
-

Registering Interrupts

- ❑ The ISR routines associated with an IRQ can be set to be either interruptible or non-interruptible when requesting the IRQ.
 - This can be conveyed through the `irqflags`.
 - ❑ `dev_name` is a historical field that is still used so system can show what IRQs are in use by what device.
-

Sharing IRQ's

- ❑ Unfortunately, there are a limited number of IRQ's in a machine.
 - ❑ So how does the operating system provide a less limited number of IRQ's?
 - ❑ It shares the same IRQ between different devices.
 - *RESTRICTION*: the device must support interrogation by the processor.
-

Sharing IRQ's (cont.)

- ❑ To share an IRQ, both device drivers when registering must indicate through the `irqflags` that they are capable of sharing.
 - If so, for the IRQ entry, a chain is built of ISRs for each device that is sharing the IRQ.
 - ❑ When the IRQ occurs, each ISR in the chain is invoked.
 - ❑ Each ISR then interrogates its hardware device to determine if it generated the interrupt.
-

Dead Hardware?

- ❑ There is a possibility that a device can not respond when expected by the device driver.
 - ❑ Process requests to read floppy disk, device driver issues read operation, but the device never responds.
 - i.e. No floppy/CD in drive
 - ❑ The operating system deals with this problem through the use of timers.
-

Timers

- ❑ Before the device driver releases control, it registers a **timer**.
 - As discussed previously...
 - ❑ If no interrupt from the device is received, the **timer** will expire and a registered function of the devices choice will be executed.
 - ❑ If the interrupt occurs then the ISR must either restart or delete the **timer** to prevent over service.
-