

# Operating Systems II

---

## Interprocess Communication

# IPC Support

---

- ☐ IPC is used to coordinate and share information between tasks.
  - ☐ Resource sharing
  - ☐ Synchronization
  - ☐ Data exchange (connectionless and connection oriented)
    - This can be in the form of one-to-one or one-to-many.
-

# Kernel Synchronization

---

- ❑ The kernel manages system resources and access to these resources must be synchronized.
  - ❑ To provide the synchronization the kernel controls who can run during a system call.
  - ❑ Other processes get “scheduled” only in three different cases...
-

# Kernel Synchronization (cont.)

---

1. When the system call invokes the `schedule()` method.
  2. When the system call invokes a method that will *suspend* the process.
  3. If pre-emptive scheduling is in force, when an interrupt occurs that affects scheduling.
-

# Kernel Synchronization (cont.)

---

- So, the kernel provides synchronization internally by:
    - Only calling `schedule()` when it will not affect synchronization.
    - Same for a suspending function.
    - Turn off interrupts when performing critical regions.
      - Simple, but does slow down the system...
-

# Multiprocessor Systems

---

- ❑ Unfortunately, the approach of enabling/disabling interrupts does not work for multiprocessor systems.
  - ❑ So how does it perform synchronization?
    - Using a spinlock (mutex).
-

# Spinlocks

---

- ❑ `typedef struct {int lock;} spinlock;`
  - ❑ Just like a regular mutex, unlocked it contains a 1, locked it contains a 0.
  - ❑ System calls try to acquire the lock by decrementing the value. Release the lock by incrementing the value.
  - ❑ Requires processor test&set atomic operation.
-

# Spinning the Tires...

---

- ❑ Unlike an OS mutex/semaphore these locks have a busy loop!!!
  - ❑ In the kernel, you cannot put the system call or ISR into a waiting state. So you just spin the loop waiting for it to be unlocked.
  - ❑ No problem ... the other processor is expected to release the lock.
  - ❑ Defined to be empty on single processor systems.
-



# Assessment of Spinlocks

---

- ☐ Can be used in interrupt service routines on multiprocessor systems.
  - ☐ Quick
  - ☐ ... and dirty. Wasted clock cycles of execution.
-

# Read/Write Locks

---

- ❑ Same as spinlocks, but ...
    - Allow multiple readers of a resource.
    - Only one writer with no readers.
  - ❑ `typedef struct {int lock; }rwlock;`
  - ❑ Lock field starts with `RW_LOCK_BIAS`
    - For example `0x01000000`
-

## Read/Write Locks (cont.)

---

- ❑ Readers try to decrement the lock and the result being a positive number.
  - ❑ Writers try to decrement the lock by `RW_LOCK_BIAS` and get a 0.
  - ❑ Also carried out by a busy loop.
  - ❑ More costly to perform than spinlocks
    - Use them more wisely!
-

# Spinlock Macros

---

- ❑ There are several macros one can provide to assist in using spinlocks & read/write locks.
  - ❑ `SPIN_LOCK_UNLOCKED` & `RW_LOCK_UNLOCKED`
    - Initialization of a spinlock or read/write lock.
  - ❑ `spin_lock()` & `spin_unlock()`
    - Lock and unlock without any affect on interrupts.
  - ❑ `read_lock()` & `read_unlock()`
  - ❑ `write_lock()` & `write_unlock()`
    - Lock and unlock without any affect on interrupts.
-

# Spinlock Macros (cont.)

---

- ❑ `spin_lock_irq( )` & `spin_unlock_irq( )`
    - Lock/Unlock, but also allow interrupt manipulation for current processor. Sets interrupt mask.
  - ❑ `spin_lock_irqsave( )` & `spin_unlock_irqrestore( )`
    - Lock/Unlock, but also allow interrupt manipulation for current processor. Restores interrupt mask.
-

# Task synchronization

---

- ❑ Application processes (tasks) are synchronized by the use of `wait_queues`.
  - ❑ As previously discussed, queues are often implemented as doubly linked lists.
-

# Queue macros

---

- ❑ Queue management only!!!
  - ❑ `DECLARE_WAITQUEUE( )`
    - Declares and initializes a `wait_queue_t` structure.
  - ❑ `add_wait_queue( )`
    - Adds the task to the queue
  - ❑ `remove_wait_queue( )`
    - Removes the task from the queue
  - ❑ `wait_queue_active( )`
    - Anything in the queue?
-

# Synchronization Methods

---

- ❑ `sleep_on( )`
    - eternal sleep in the queue.
  - ❑ `interruptible_sleep_on( )`
    - sleep, but can be awoken by an interrupt.
  - ❑ `sleep_on_timeout( )`
    - sleep for a fixed amount of time.
  - ❑ `interruptible_sleep_on_timeout( )`
    - sleep for a fixed amount of time AND can be interrupted from sleep.
-



# Synchronization Methods (cont.)

---

- ❑ `wake_up( )`
    - Wake the task from the queue.
  - ❑ `wake_up_nr( )`
    - Wake a given number of tasks from the queue
  - ❑ `wake_up_sync( )`
    - Wake up tasks, but do not schedule until next **regular** scheduling period.
  - ❑ `interruptible wake_up( )` variants
    - Wakes up ONLY tasks that were put to sleep with interruptible status.
-

## sleep\_on( )

---

```
Sleep_on(struct wait_queue **p) {  
    struct wait_queue wait;  
    current->state = TASK_UNINTERRUPTIBLE;  
    wait.task = current;  
    add_wait_queue(p, &wait);  
    schedule();  
    remove_wait_queue(p, &wait);  
}
```

---

# Semaphores

---

```
Struct semaphore {  
    int count;  
    int sleepers;  
    wait_queue_head_t *wait;  
}
```

- ❑ count - The lock! Atomic for test&set in one operation. Supported by hardware architecture.
- ❑ sleepers –  $\text{sleepers} + \text{count} = \text{correct value for semaphore}$ .
- ❑ wait - The list of tasks that are sleeping.

## Semaphores (cont.)

---

- ❑ `up( )` – increment count and wake ALL tasks if  $\leq 0$ ; wake a process if count = 1.
  - ❑ `down( )` – decrement count; modify sleepers; and sleep if count  $< 0$ .
  - ❑ `down_interruptible( )` – task is interruptible if it does go to sleep on trying to lock.
  - ❑ `down_trylock( )` – does not block task if lock cannot be acquired.
-

# Communication via Files

---

- ❑ Oldest mechanism of exchanging data.
  - ❑ In a multitasking environment you require synchronization to ensure correctness with the file.
  - ❑ Most kernels provide file locking...
    - You can lock the full file or just areas of the file.
    - You can have *mandatory* or *advisory* locking.
-

# Implementation Details

---

- ❑ The kernel maintains a linked list of `file_lock` structures for each lock.
  - ❑ The list is per file!
  - ❑ Different information stored in the structure include
    - next pointer, list of waiting processes, file indicator, start & end of lock, lock type, ...
-

# File Control

---

- ❑ `sys_fcntl(int fd, int cmd, void* arg) ;`
  - ❑ `fd` – file descriptor
  - ❑ `cmd` – command to perform on the file
  - ❑ `arg` – argument(s) needed to perform the command.
-

# Locking Files

---

- ❑ cmd can be one of F\_GETLK, F\_SETLK, or F\_SETLKW.
- ❑ arg – flock structure

```
Struct flock {  
    short l_type; /* F_RDLCK, F_WRLCK, F_UNLCK */  
    short l_whence; // SEEK_SET, SEEK_CUR, SEEK_END  
    off_t l_start; // offset relative to l_whence  
    off_t l_len; /* length of area to lock */  
    pid_t l_pid; /* returned with F_GETLK */  
}
```

---



## Locking Files (cont.)

---

- ❑ `F_GETLK` – tests whether the lock is possible, if not attempted lock returned.
  - ❑ `F_SETLK` – sets the lock specified, returns either way.
  - ❑ `F_SETLKW` – sets the lock specified, but blocks if lock cannot be set.
-

# Pipes

---

- ❑ Pipes are used to share data between two distinct processes.
  - ❑ The pipe uses a go between `inode` (file system block) to store the information.
  - ❑ Once created, processes use `read()` and `write()` to interact with the pipe.
-

## Pipes (cont.)

---

- ❑ Creating a pipe involves
    - allocating an `inode` for storing the data.
    - generating a reader and writer file descriptor for the `inode`.
  - ❑ Accessing the pipe is the same as for a file.
-

# pipe\_inode\_info structure

---

```
wait_queue_head_t wait; /* wait queue for processes */
char *base; /* address of buffer */
int start; // amount written and yet to be read
int len; /* current amount that is unread. */
int readers; /* # of processes reading */
int writers; /* # of processes writing */
int waiting_readers; // # blocked readers in Q
int waiting_writers; // # blocked writers in Q
int r_counter; // # read processes that opened
int w_counter; // # write processes that opened
```

---