

Operating Systems II

Process Management

Process Management

- We already talked about the `task_struct` data structure that holds all of the information about the task.
 - Also talked about
 - the states a task can exist in.
 - the queues the task can be held in.
 - What about the other aspects of process management.
-

Task Creation

- ❑ Every task is created from another task.
 - Maybe from the system idle task?
 - ❑ A new task is created when a parent process invokes either a `fork()` or `clone()` system call.
 - `clone()` – inherits more from the parent task.
 - ❑ The system calls eventually call the `do_fork()` function that creates a new task.
-

`do_fork()`

- ☐ Allocate a new `task_struct` data structure.
 - ☐ Link the `task_struct` into the process table.
 - ☐ Create a new kernel space stack for execution when inside the kernel.
 - ☐ Copy the fields from the parent's `task_struct` into the child's.
-

`do_fork()` (cont.)

- ❑ Modify the child's `task_struct` fields specific to itself.
 - New process identifier
 - Links to the tasks parent and siblings
 - initialize the process specific timers.
(creation, time quantum, and so on)
-

do_fork() (cont.)

- ❑ Copy other data structures that are in parent and should be replicated for the child task.
 - File table and new file descriptor for each open file.
 - Create a new user data segment and copy the data.
 - Copy signal and signal handling information.
 - Copy virtual memory tables.
 - ❑ Change the child's state to RUNNING.
-

Task Termination

- ❑ Several ways to terminate a task
 - task making the `exit()` system call.
 - delivered a signal with the signal handler disposition to die.
 - Being forced to die under certain exceptions.
 - ❑ Termination work is done by `do_exit()`.
 - ❑ What does `do_exit()` do?
-

`do_exit()`

- ☐ Sets a global kernel lock.
 - ☐ Sets the task state to ZOMBIE
 - ☐ Notifies any child of process termination.
 - ☐ Notifies the parent with an exit signal (usually SIGCHILD).
 - ☐ Releases any resources allocated by `do_fork()` (such as open files, ...)
 - ☐ Calls `schedule()` (never returns).
-

Task Scheduling

- Tasks are scheduled via the `schedule()` function.
 - What does `schedule()` do?
 - Exactly what happens is dependent on the scheduling algorithm.
-

Schedule () - Overview

- ❑ Performs various periodic work (such as running the bottom halves of interrupt handlers)
 - ❑ Inspect the set of tasks in the running state.
 - ❑ Chose one task to execute based on the scheduling policy.
 - ❑ Dispatch the task to run on the CPU until an interrupt occurs.
-

Schedule ()

- ❑ Release the global kernel lock.
 - ❑ Do any software interrupts (2nd stage).
 - ❑ Grab current process and current CPU.
 - ❑ Set state appropriately.
 - ❑ Find next process to schedule.
 - ❑ Use `switch_to()` macro to perform the transfer (saves state of old task, loads state of new task).
 - Architecture Specific!!!
-

Process Table Support

- ❑ Some common functions found in a kernel to assist process management
 - `init_task` – The first task in the process table.
 - `find_task_by_pid(x)` – given a pid will return the `task_struct` pointer.
 - `for_each_task(p)` – For loop construct that will allow you to iterate over each task in the process table.
-