

Operating Systems II

System Calls

Motivation of System Calls

- ❑ System calls describe the interface that exists between user programs and the kernel.
 - ❑ This interface has to be flexible & powerful
 - Allows for easy development of applications.
 - ❑ The interface also has to be clear and controlled
 - Necessary for *any* security support the kernel wants to provide.
-

System Calls

- ❑ Define a transition from user mode to system mode.
 - ❑ How the transition is made can vary dramatically.
 - ❑ Key factors include:
 - Kernel design – Kernel thread to offer all services.
 - Processor architecture – Processor modes and addressing support.
-

Simple Design

- ❑ The simplest approach to system calls is to provide interrupt “wrapping”.
 - ❑ That is, each of your system calls starts with disabling interrupts and ends with enabling interrupts.
 - ❑ This results in the system call taking control of the system without any possibility of pre-emption.
-

Problems

- ❑ No support for kernel developer.
 - Careful to ensure all system calls deal with interrupts properly.
 - ❑ The approach works but makes some vital assumptions!
 - Disabling interrupts will not cause application problems.
 - Kernel trusts the user!
-

Alternative Approach

- ❑ An alternative approach to system call is to invoke ALL system calls through a common interface.
 - ❑ This interface will perform the transition between system and user modes.
 - ❑ A common interface is through an interrupt!
 - Linux uses interrupt 0x80!
-

What System Call?

- ❑ Each system call has a registered number.
 - ❑ When the system call is made, the *user mode* process will
 - put the system call id into a register
 - push parameters into registers (or onto a stack)
 - then generate the interrupt!
 - ❑ The kernel interrupt service routine then takes over.
-

system_call ISR

- The global system call handler.
 - saves the registers
 - restricts memory accessing to kernel space.
 - grabs system call id and verifies it is legit.
 - checks if task is being traced/profiled.
 - invokes the system call – with parameters.
-

system_call ISR (cont.)

- ☐ After the specific system call function the handler will...
 - Place result of system call in register.
 - Perform system maintenance (scheduling of tasks, check for pending signals, ...)
 - ☐ Why not save time since we are in the kernel anyway?
 - Return from interrupt.
-

System Call Development

- ❑ Writing a system is simpler for the kernel developer
 - Write your system call function
 - Pick a system call id
 - Add function pointer and system call id to kernel table.
 - ❑ Some limitations:
 - Finite number of system calls
 - Up to a fixed number of parameters
 - ❑ Remember as well functionality in the system call is limited!
 - No access to shared libraries
-

System Call Stubs

- ❑ A system call stub is generated to aid the user application developer.
 - They have to load registers with parameters and system call id then generate the system call interrupt!
 - ❑ The stub can be generated by the kernel during compilation.
 - `myKernel_foo()`
-

Where's the Beef in Linux?

- `sys_call` – routine to handle interrupt 0x80
 - `arch/i386/kernel/entry.s`
 - `sys_call_table` – table of registered system calls.
 - `arch/i386/kernel/entry.s`
 - `NR_syscalls` – number of system calls registered (hard coded)
 - `arch/i386/kernel/entry.s`
-

Creating a Linux System Call

- Write your system call function (kernel/sys.c)
 - Need to use `asm linkage` in declaration
 - Give your system call an id number
 - Register system call in (order matters)
`arch/i386/kernel/entry.s`
 - Make sure to increment `NR_syscalls` (same file)
 - Add your system call to the `sys_call_table` (`include/asm/unistd.h`)
 - `#define __NR_my_sys_call 243`
-

Creating a Linux System Call

- Specify the new system call in the user application.
 - `_syscall#` facility to generate the stub.
 - Remember, up to 6 arguments supported!
-
- Or you can invoke the system call without a stub –
must `#include <sys/syscall.h>`
 - `syscall(194, &baz, bar);`
-